

# VMware {code} Experience Training

Open Source Projects for Enterprise Cloud Native

Ben Corrie, Senior Staff Engineer

Henry Zhang, Chief Architect, R&D, VMware China

Patrick Daigle, Sr. Tech Marketing Architect, VIC

Steven Zou, Staff Engineer, Project Harbor



© 2014 VMware Inc. All rights reserved.


# Disclaimer

- This presentation may contain product features that are currently under development.
- This overview of new technology represents no commitment from VMware to deliver these features in any generally available product.
- Features are subject to change, and must not be included in contracts, purchase orders, or sales agreements of any kind.
- Technical feasibility and market demand will affect final delivery.
- Pricing and packaging for any new technologies or features discussed or presented have not been determined.

# VMware's Participation in Open Source Community

**Use** open source code within the enterprise:

- For implementation of industry standards
- To improve software usability
- To accelerate our software development processes



**Contribute** to upstream projects to add features or fix bugs:

- Linux kernel drivers for our virtual hardware, but also work on key kernel infrastructure
- Open Stack improvement for network, storage, stability
- Kubernetes integration with vSphere as compute and storage provider
- Many smaller contributions and bug fixes to a wide variety of projects

**Create** and release a new open source project and build new community to support it

- Variety of project – from utilities and glue code to significant, standard-setting projects

# Leadership in Open Source Cloud Native Projects



vSphere Integrated  
Containers Engine



HARBOR™

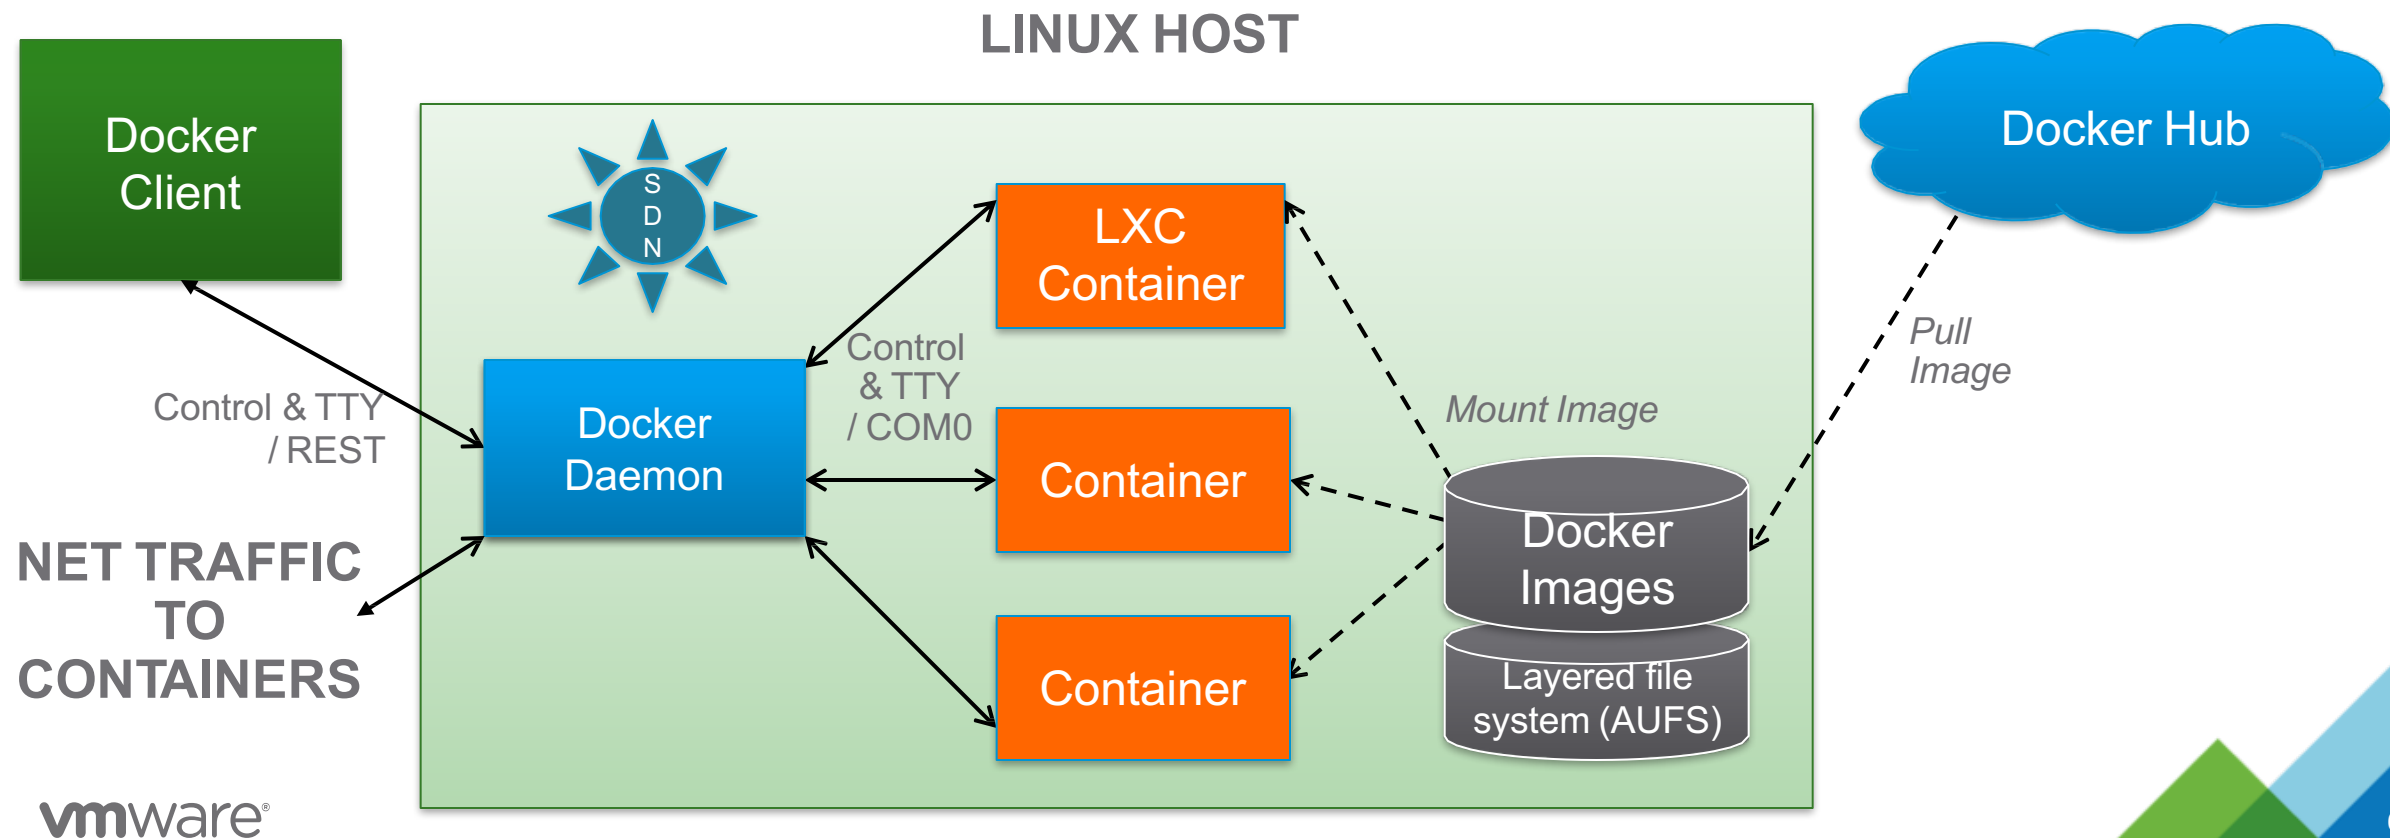


ADMIRAL™

# vSphere Integrated Containers

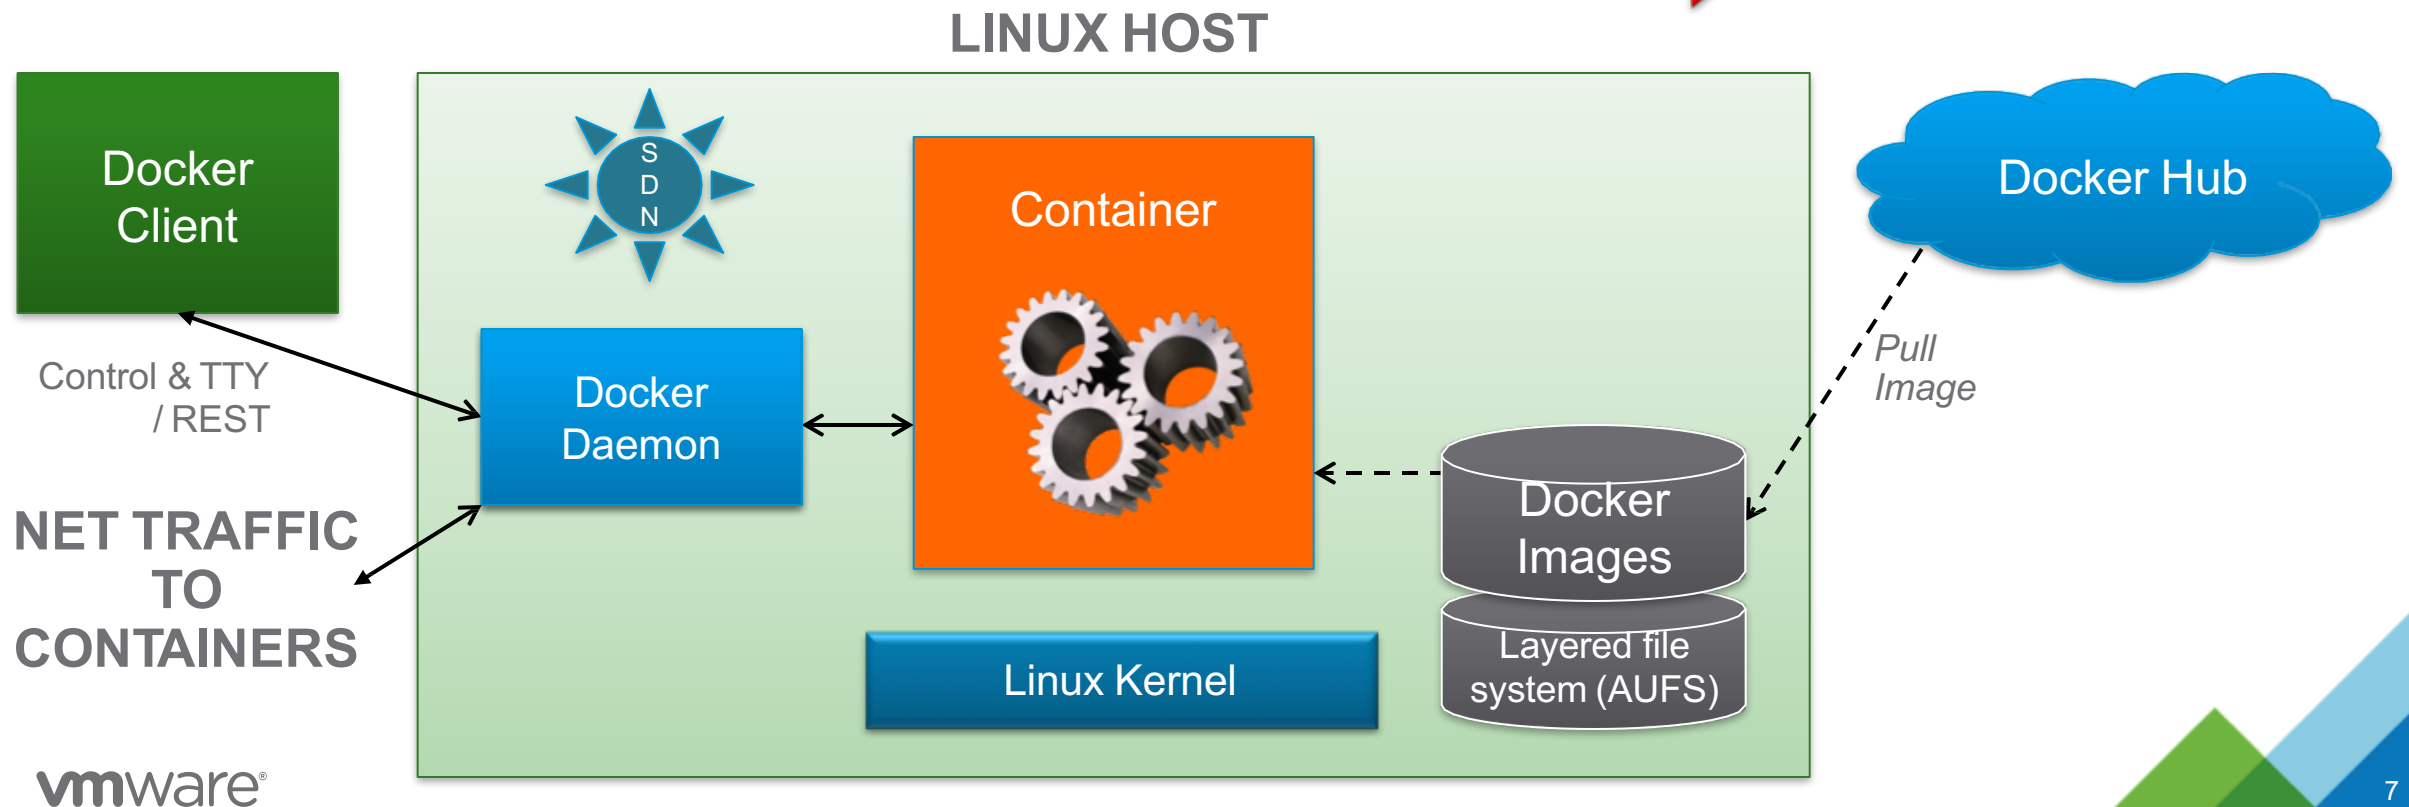
# Docker 101

```
vmware [ ~ ]$ sudo docker run -it ubuntu
root@636a8cb6d180:/# ls
bin    dev    home  lib64  mnt    proc   run    srv    tmp    var
boot  etc    lib   media  opt    root   sbin   sys    usr
root@636a8cb6d180:/#
```



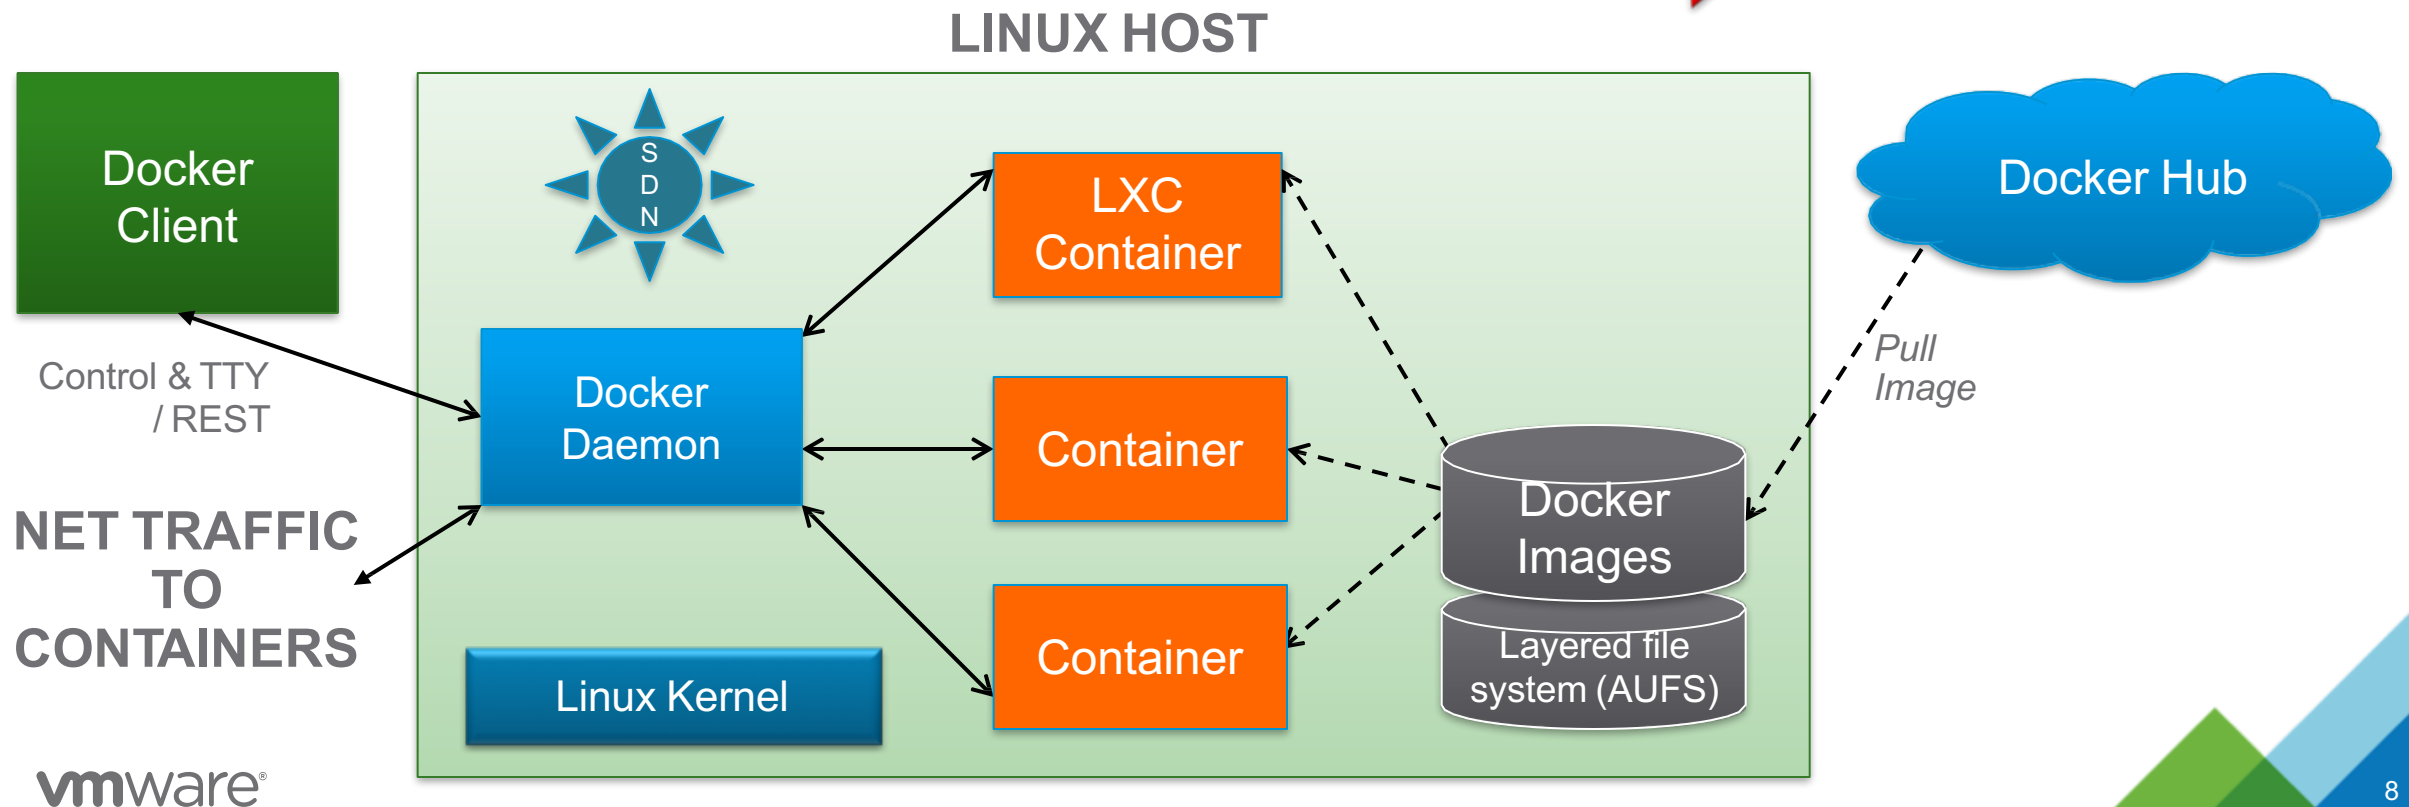
# What is a Container?

1. An executable process
2. Resource constraints / private namespace
3. Binary dependencies: Application, runtime, OS
4. A shared Linux kernel for running the executable
5. Ephemeral and persistent storage layers



# What is a Container Host?

1. Control plane & lifecycle management for containers
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A single Linux kernel manages everything
5. A static size and a resource reservation when virtual






# What is a Hypervisor?

1. Control plane & lifecycle management for VMs
2. Resource scheduling and a VM abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A hypervisor kernel manages everything except apps
5. A static size, but no resource concerns unless nested

STATEFUL  
LONG-RUNNING  
MULTI USER  
MULTI USE

## ESX HOST



# Types of Containers

- ***Long-running***
  - Can be stateless or stateful
  - Eg. Application servers, databases, load-balancers, KV stores etc
  - Typically a need for strong isolation
- ***Transactional***
  - Runs for a period and transforms some data
  - Eg. Runs a build. Processes a web request. Batch processing
  - Should only consume resource when running
- ***Sidecar / helpers***
  - Augments the capabilities of a service or provides a helper function
  - Eg. Logging, monitoring, caching
  - Scales with the service. Potentially hindered by strong isolation

# Isolation Domains and Data Persistence

- How do I isolate workloads from each other?
  - Runtime isolation – resource limits, kernel panic, ESX host failure, rack or region failure
  - Network isolation – traffic sniffing, firewalls, encryption, rate limiting
  - Storage isolation – data persistence, backup, networking, RBAC
- Stateful vs Stateless / Cattle vs Pets
  - Different classifications of data. Where should it go?
    - Image state, container state, volume state. What's the difference?
  - Should data lifespan be inherently tied to compute (VM / container)?
  - Without live migration, shared storage and HA, stateless looks attractive
- The question is not “what plumbing do I need?”
  - What *characteristics* or business value do I need for my application?
  - Better question than, “do I need a container or a VM?”

# Characteristics: Containers vs VMs? So 2014!

- **Speed**
  - Start time vs throughput. Benefits transactional containers
  - Hello World vs Tomcat. Less benefit for long-running apps
- **Efficiency**
  - Less memory consumption? Depends when virtualized
  - Network traffic NAT'd through guest vs straight to vNIC
- **Portability**
  - Docker image abstraction and API is very portable
  - Subtle issues with kernel versions and patches
- **Isolation**
  - Shared kernel significantly reduces runtime isolation
- **Granularity**
  - Containers are great for granular services. Group in a VM as isolation domain.

# Container Workflow Efficiencies

- **A Portable Runtime Platform**
  - “Package once, run anywhere”
  - Defacto runtime model for higher level frameworks
- **The Container Model**
  - “Docker is to apt what apt is to tar” – dependency management
  - Snapshotting, image format, Docker Hub, state management (volumes etc)
  - Scripting and automation – express an environment in a text file
  - Predictable initialization state – predictability vs. reconfiguration
- **Continuous Integration**
  - Secure registries, integration with popular tooling
  - Rapid improvements in tooling, monitoring, log integration etc.

# The Container Stack and Its Challenges

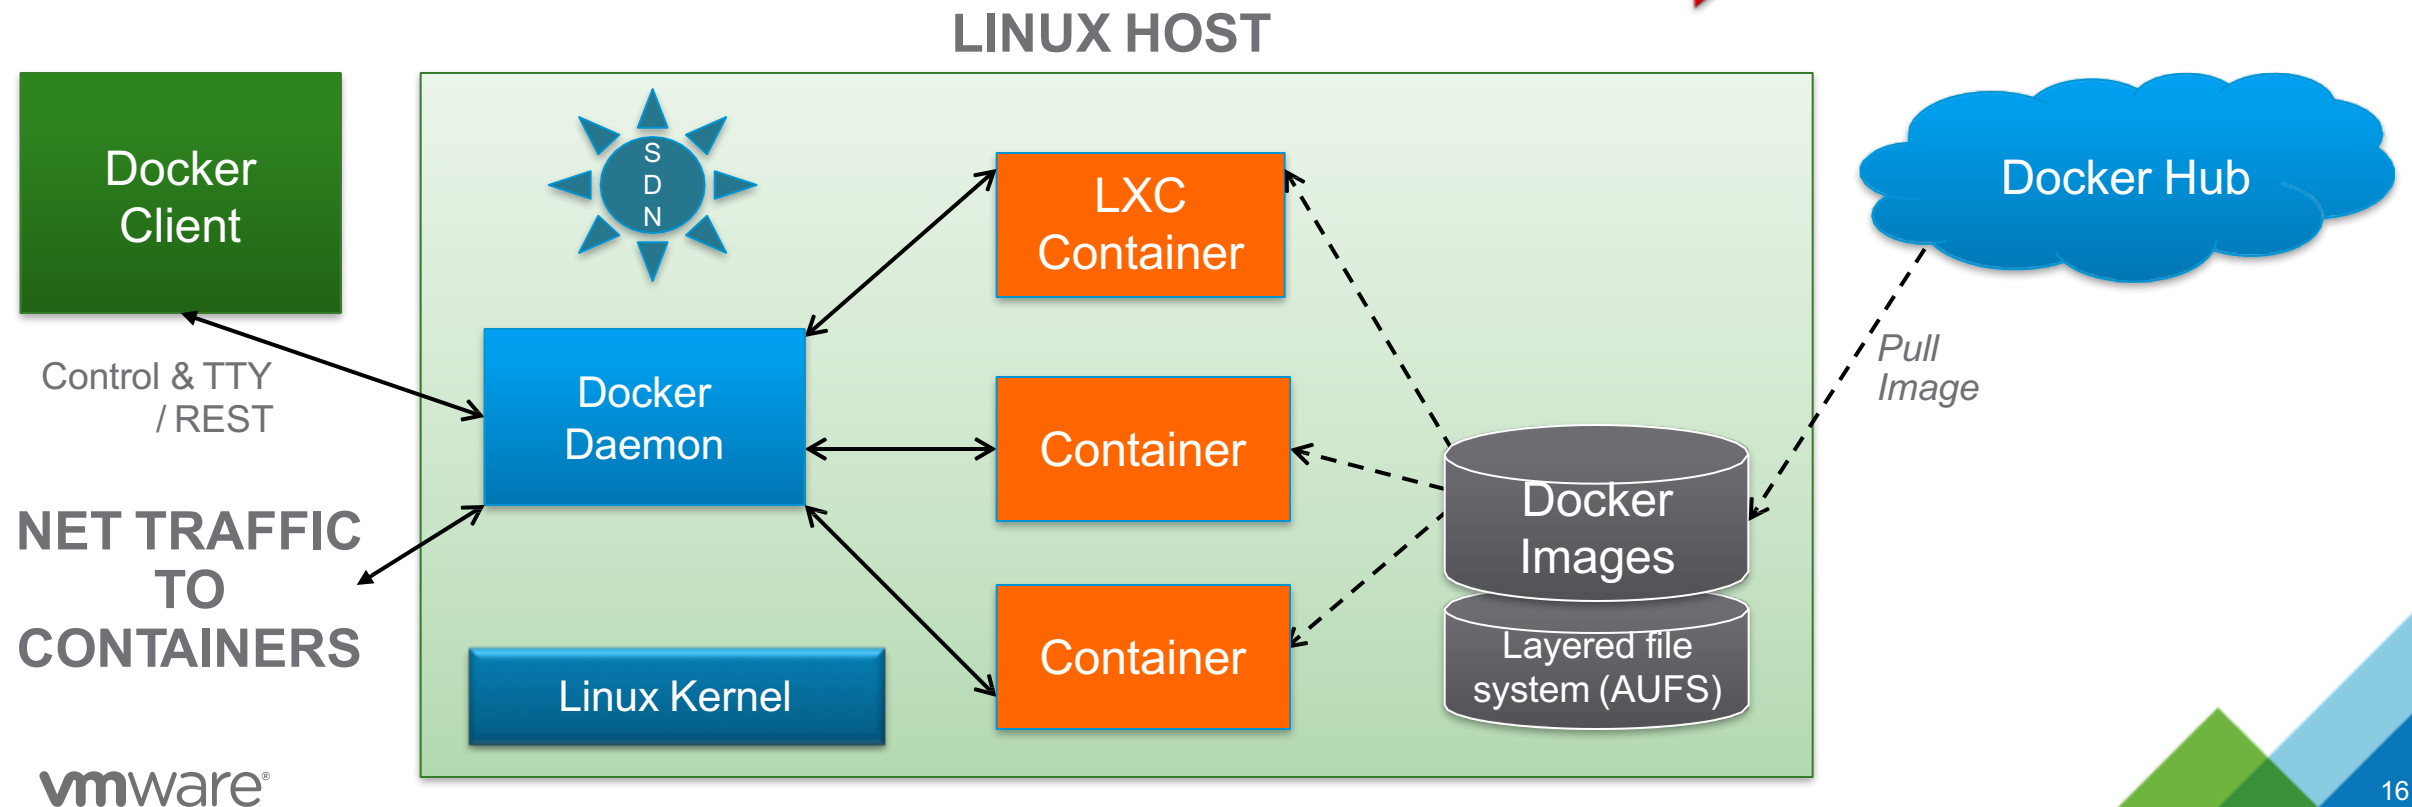
- I want to run containers. Where do I start?
- Should I run containers alongside my other apps or create a silo?
- Should I be considering bare metal?
- So how many containers should I have in each VM?
- How big should each container host be? Impact of re-configuration?
- Which containers can be safely co-located with which other containers?
- Which Linux distro do I want and how to handle the high patch cadence?
- Who is responsible for infrastructure admin and how does it integrate?

# So What is VIC and How does it help?

- **VIC brings all of the container workflow efficiencies to vSphere infrastructure**
  - Control vSphere infrastructure from a Docker client without having any vSphere credentials
  - Treat VMs as ephemerally as containers
  - No more OVAs, VMDKs, Templates, Cloning. Push / pull your state from secure registries
- **VIC allows you to translate business value into plumbing**
  - Eg. I need to deploy Wordpress with MySQL.
  - Do I want strong isolation between these workloads? Container *as* a VM
  - Do I want strong isolation from other tenants? Container *in* a VM
- **VIC helps to draw clear lines between admins and users**

# Revisit: What is a Container Host?

1. Control plane & lifecycle management for containers
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A single Linux kernel manages everything
5. A static size and a resource reservation when virtual



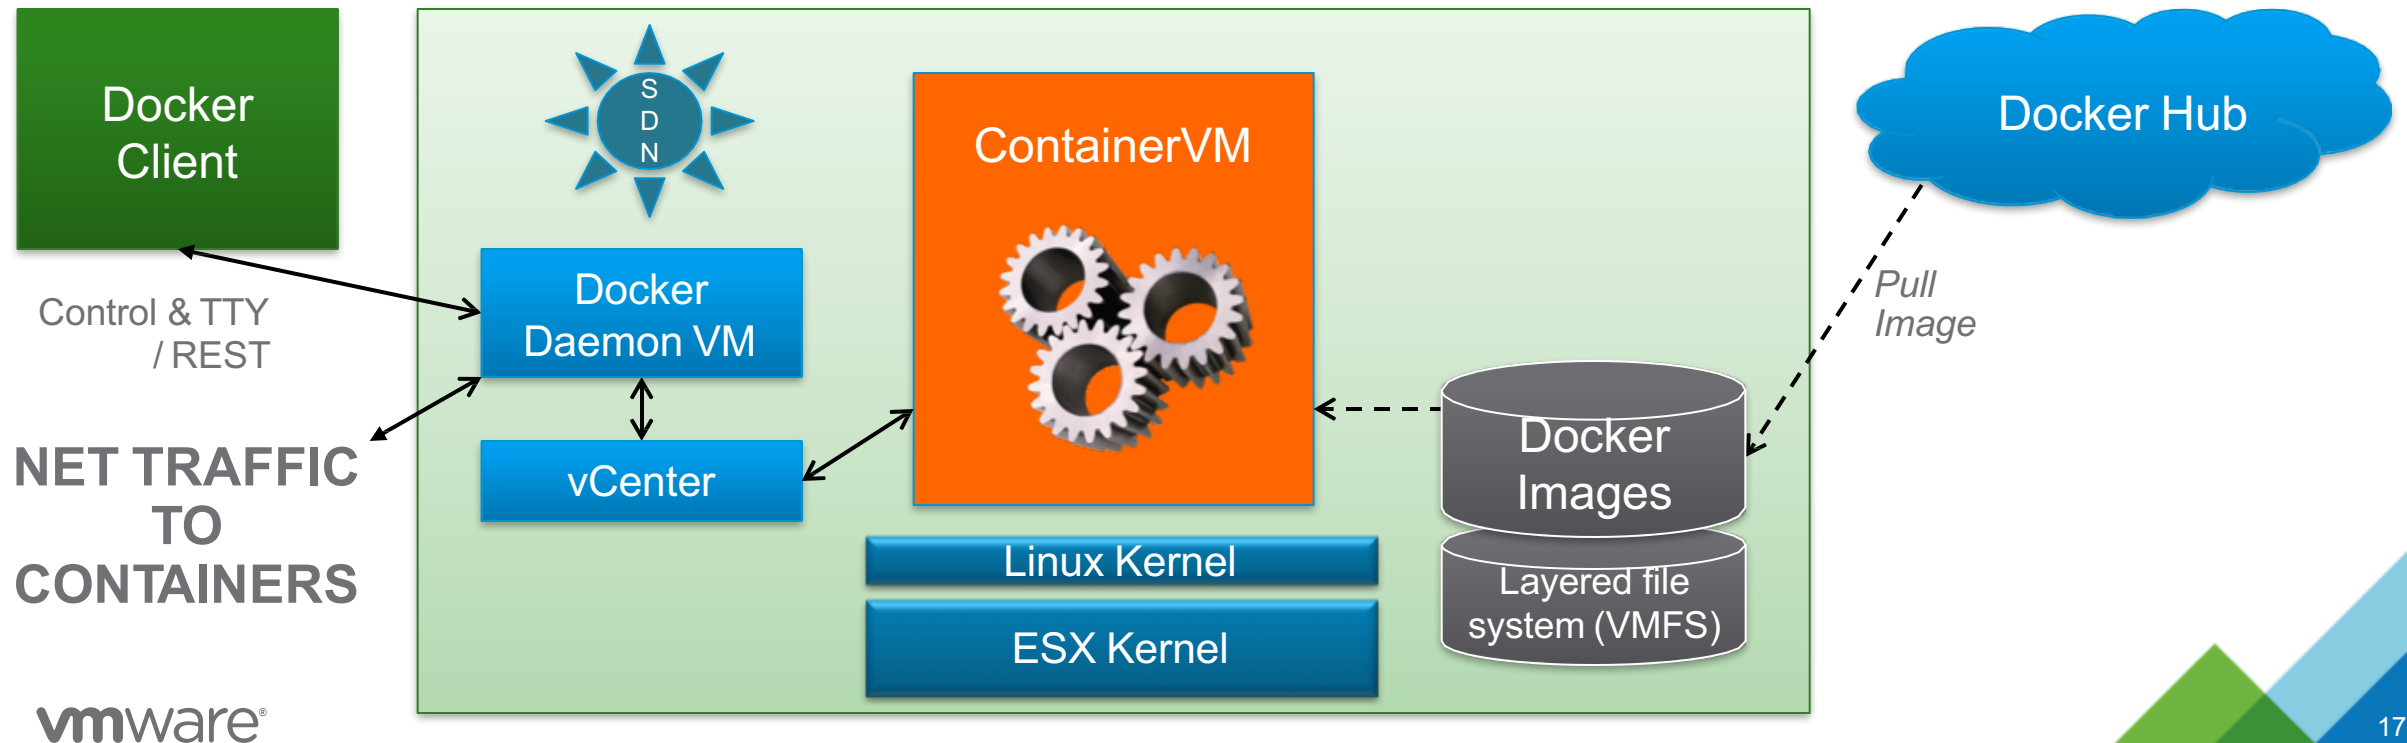


# What is a “ContainerVM”?

1. An executable process
2. Resource constraints / private namespace
3. Binary dependencies: Application, runtime, OS
4. A *private* Linux kernel for running the executable
5. Ephemeral and persistent storage layers



## ESX HOST / HYPERVISOR



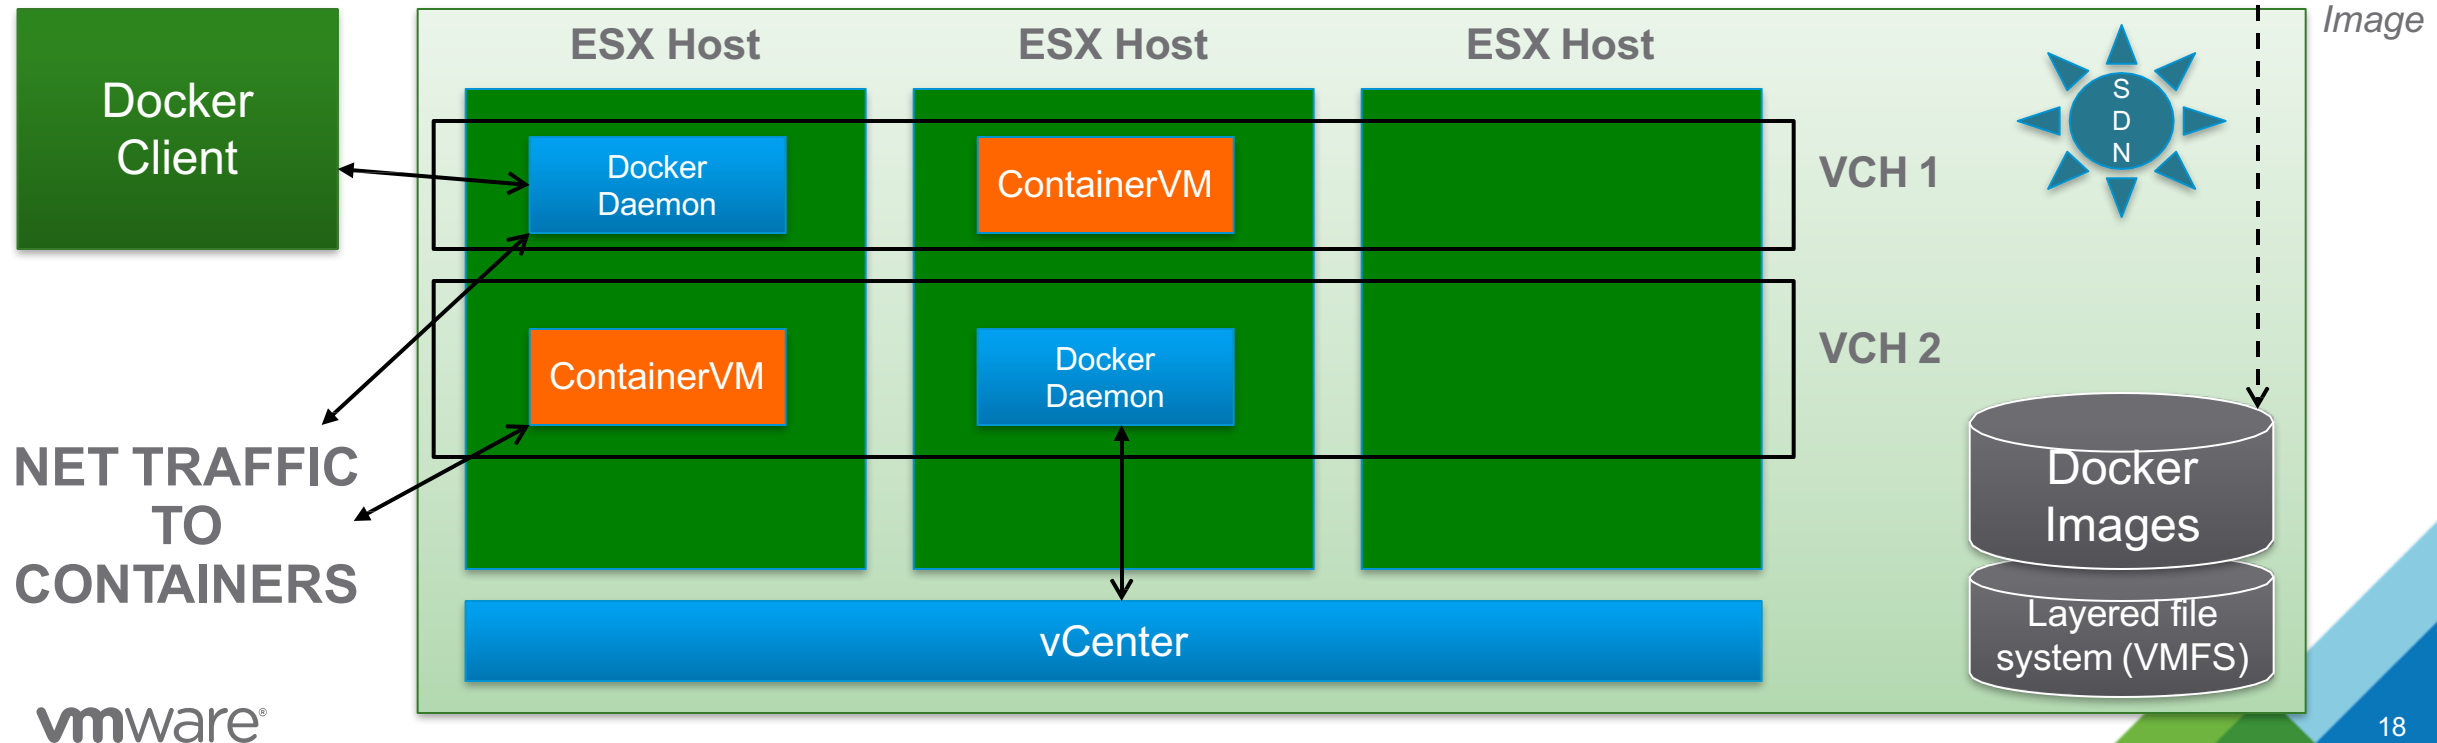
# What is a Virtual Container Host?

1. Control plane & lifecycle management for ContainerVMs
2. Resource scheduling and a container abstraction
3. Infrastructure abstractions: Storage, networking etc
4. A Linux kernel per container, separate from control plane
5. Dynamic size and a resource *limit*, not reservation!

VIRTUAL  
SINGLE USER  
MULTI USE  
MULTI HOST

Docker Hub

## vSphere Cluster



# VIC Roadmap

- **VIC 1.0**
  - Docker engine and registry (Harbor) shipped in Q4 2016
  - MVP Docker commands. Fundamentals of integration
  - vSAN, vMotion, DRS, NSX distributed port groups
- **VIC 1.1**
  - Additional Docker capabilities – Eg. compose, exec
  - Critical bug fixes
  - Deeper SDDC integrations – Eg. HA
- **VIC 1.2**
  - Security enhancements – Image signing, vulnerability scanning
  - Additional Docker capabilities
  - Deeper NSX integration
  - Further SDDC integrations – vVols, SDRS

# Questions

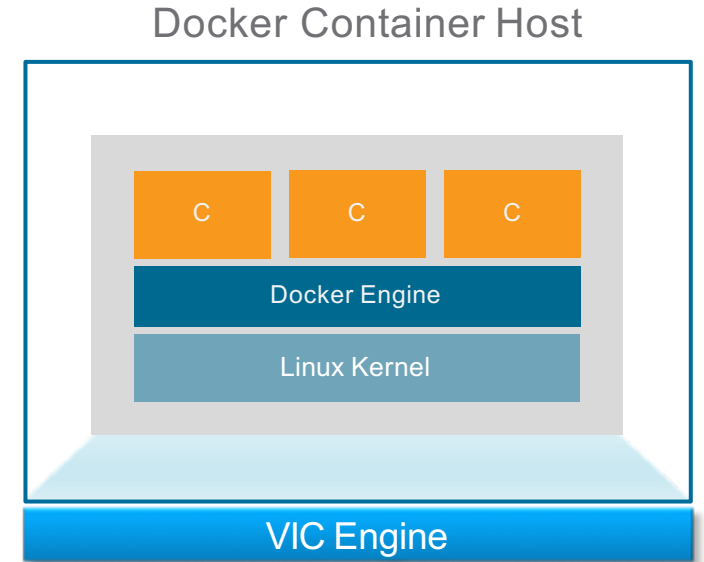
@bensdoings

[Github.com/vmware/vic-product](https://github.com/vmware/vic-product)

# Demo - VIC

## New in 1.2 - Developer Sandbox

- Developer self-service with VI Admin governance
  - Developer consumes resources via Docker API/CLI
- Provides developers with self-service for applications not yet in the enterprise service catalog
  - Rapid prototyping
- Run a full-fledged docker engine as a ContainerVM using vSphere Integrated Containers and the Docker API/CLI



```
docker run -p 12375:2375 -d vmware/dch-photon
```

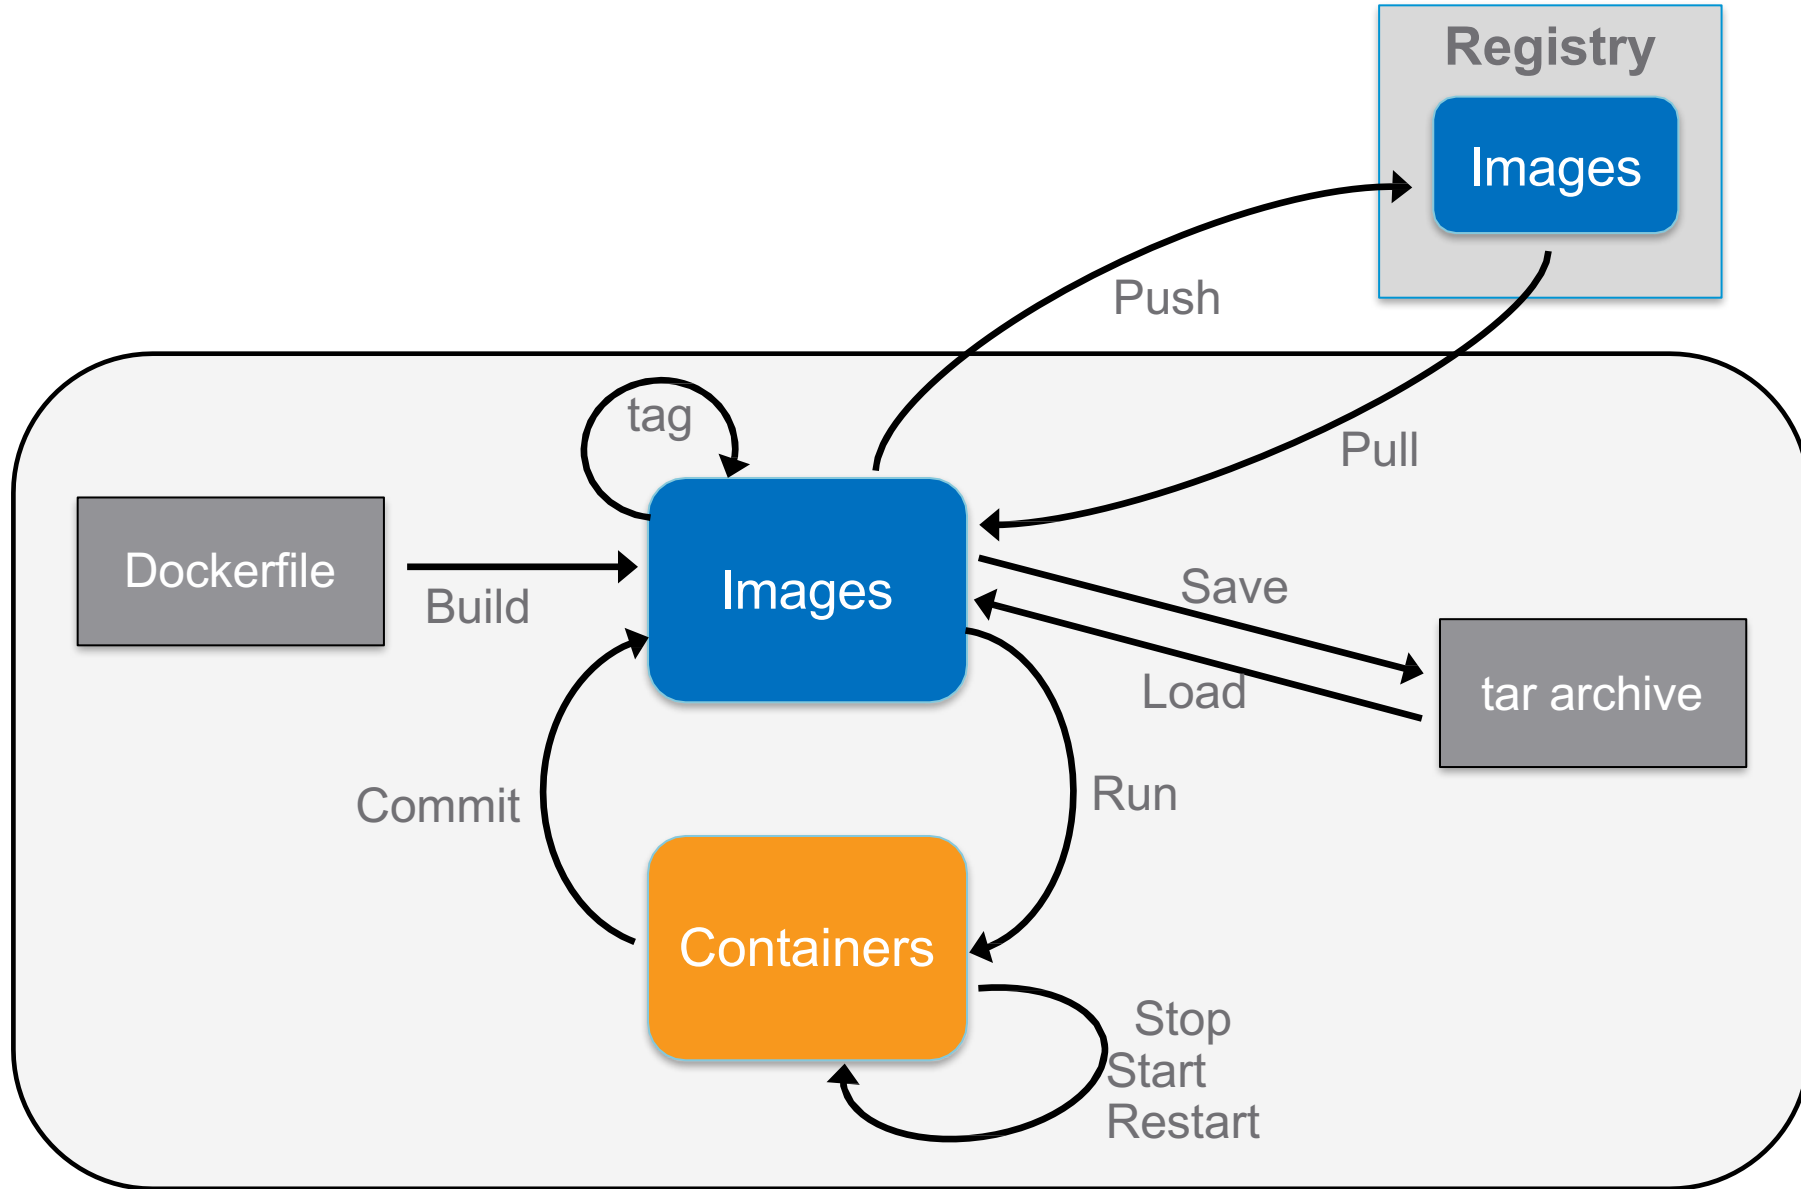
# Project Harbor

# Harbor

- 
- 1 Container Image Basics
  - 2 Project Harbor Introduction
  - 3 Consistency of Images
  - 4 Security
  - 5 Image Distribution
-

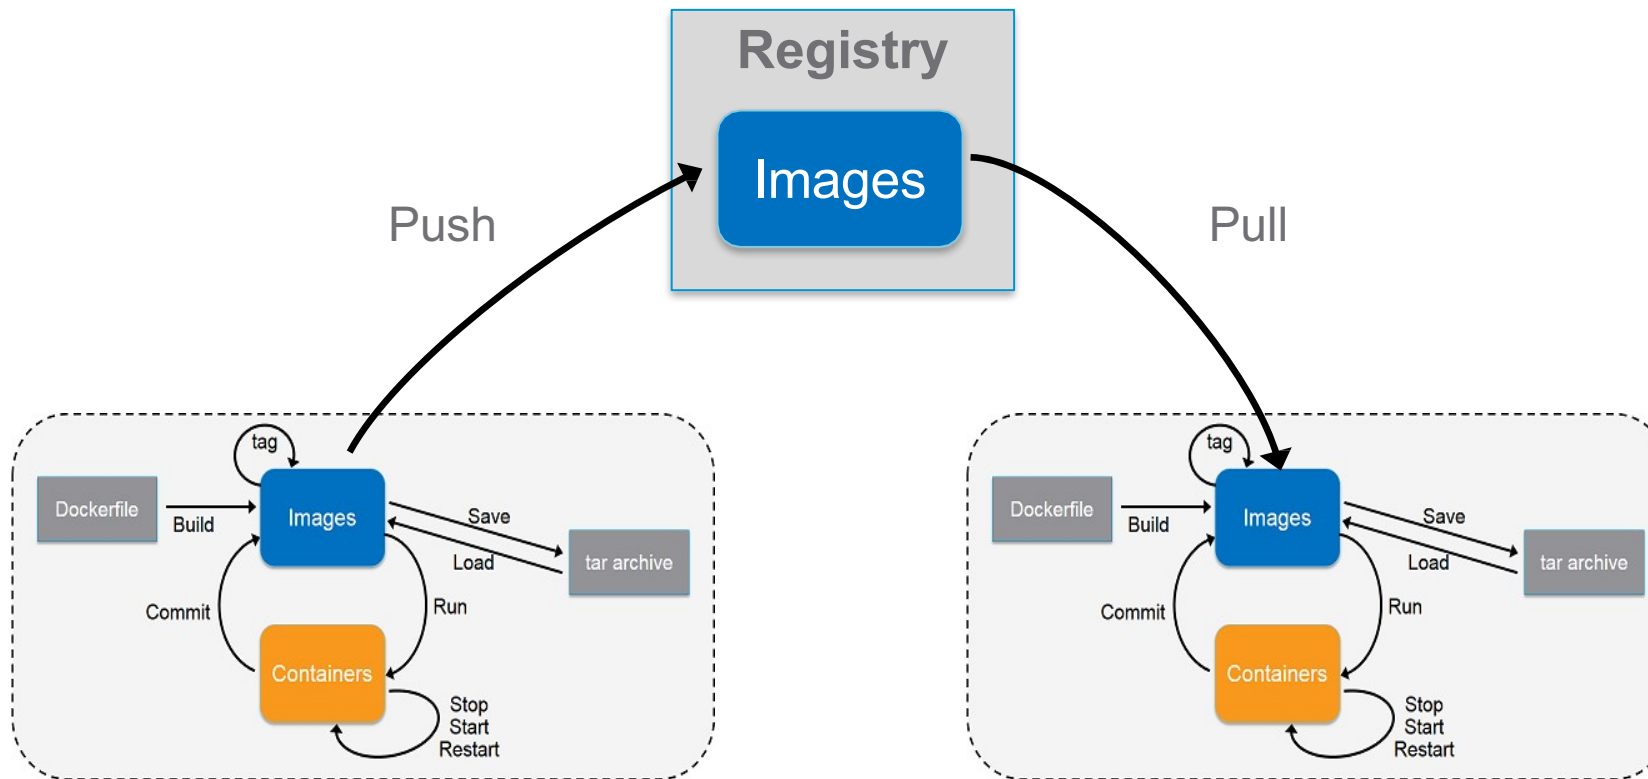


# Lifecycle of Containers and Images



# Registry - Key Component to Manage Images

- Repository for storing images
- Intermediary for shipping and distributing images
- Ideal for access control and other image management



# Harbor

- 
- 1 Container Image Basics
  - 2 [Project Harbor Introduction](#)
  - 3 Consistency of Images
  - 4 Security
  - 5 Image Distribution
  - 6 High Availability of Registry
-

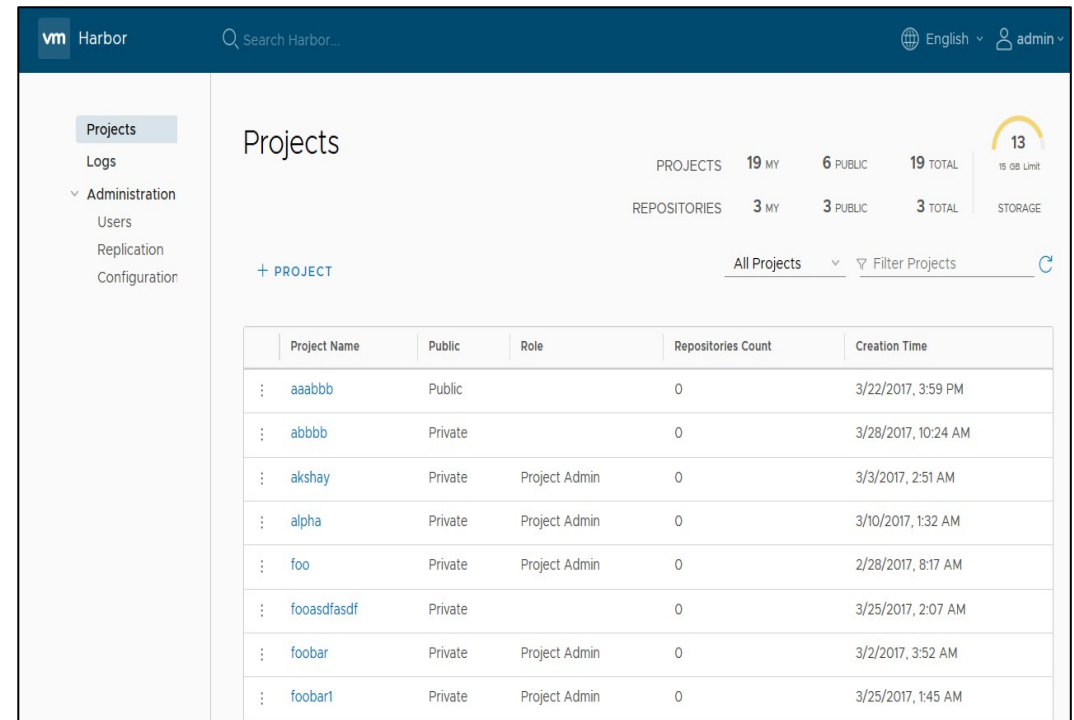
# Project Harbor



- An open source enterprise-class registry server.
- Initiated by VMware China, adopted by users worldwide.
- Integrated into vSphere Integrated Containers.
- Apache 2 license.
- <https://github.com/vmware/harbor/>

# Key Features

- User management & access control
  - RBAC: admin, developer, guest
  - AD/LDAP integration
- Policy based image replication
- Notary
- Vulnerability Scanning
- Web UI
- Audit and logs
- Restful API for integration
- Lightweight and easy deployment



Project Name	Public	Role	Repositories Count	Creation Time
aaabbb	Public		0	3/22/2017, 3:59 PM
abbbb	Private		0	3/28/2017, 10:24 AM
akshay	Private	Project Admin	0	3/3/2017, 2:51 AM
alpha	Private	Project Admin	0	3/10/2017, 1:32 AM
foo	Private	Project Admin	0	2/28/2017, 8:17 AM
fooasdfsdf	Private		0	3/25/2017, 2:07 AM
foobar	Private	Project Admin	0	3/2/2017, 3:52 AM
foobar1	Private	Project Admin	0	3/25/2017, 1:45 AM

# Users and Developers

- **Users**



Downloads



Stars



Users

- **Developers**



Forks



Contributors



Partners

## Harbor users and partners

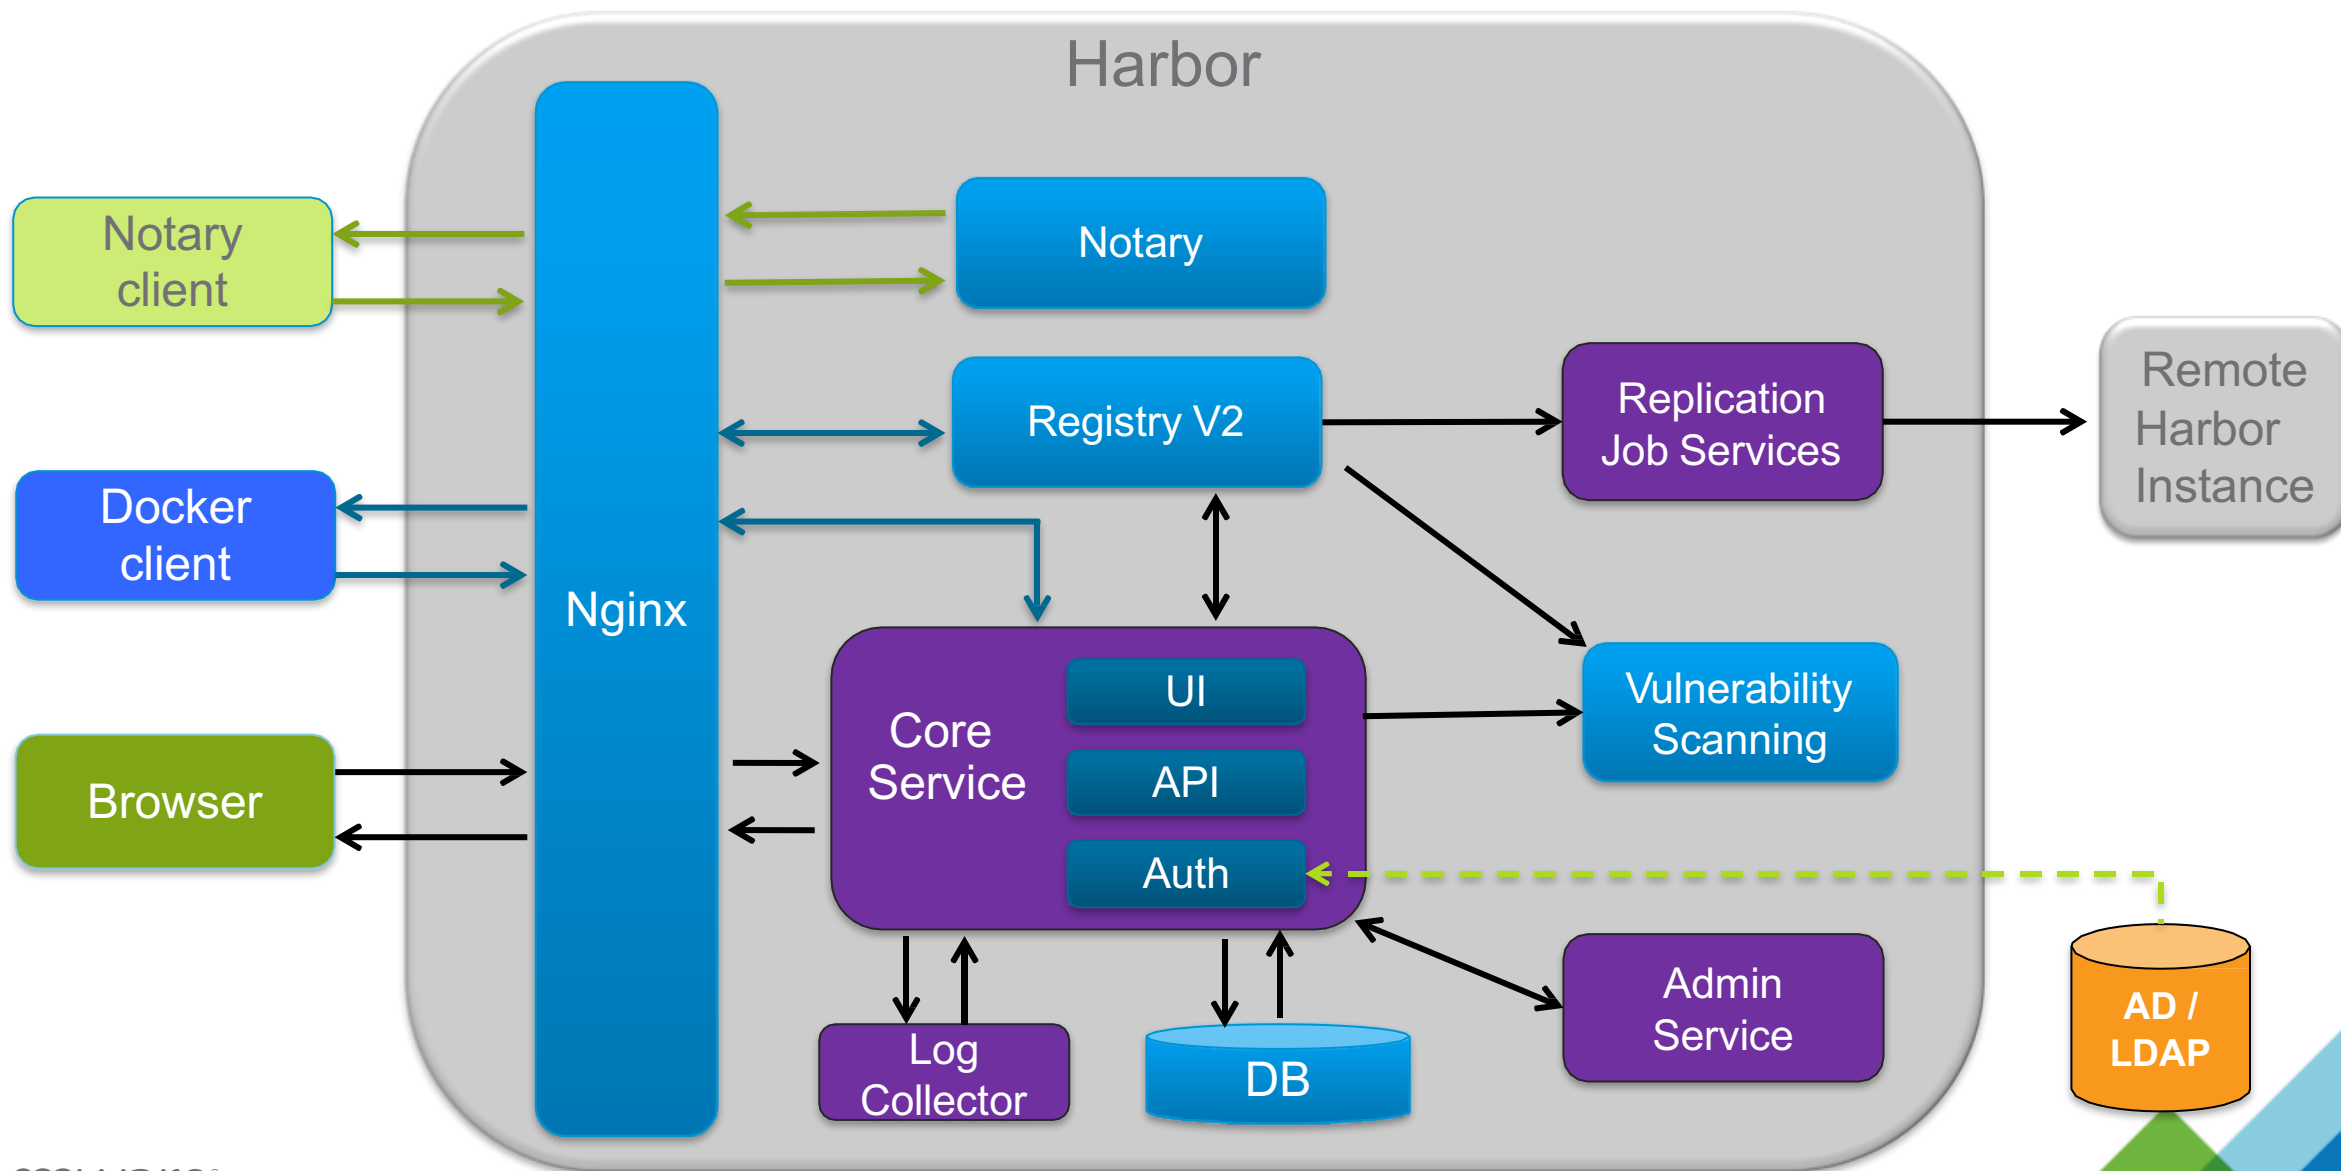


安盛天平

TalkingData

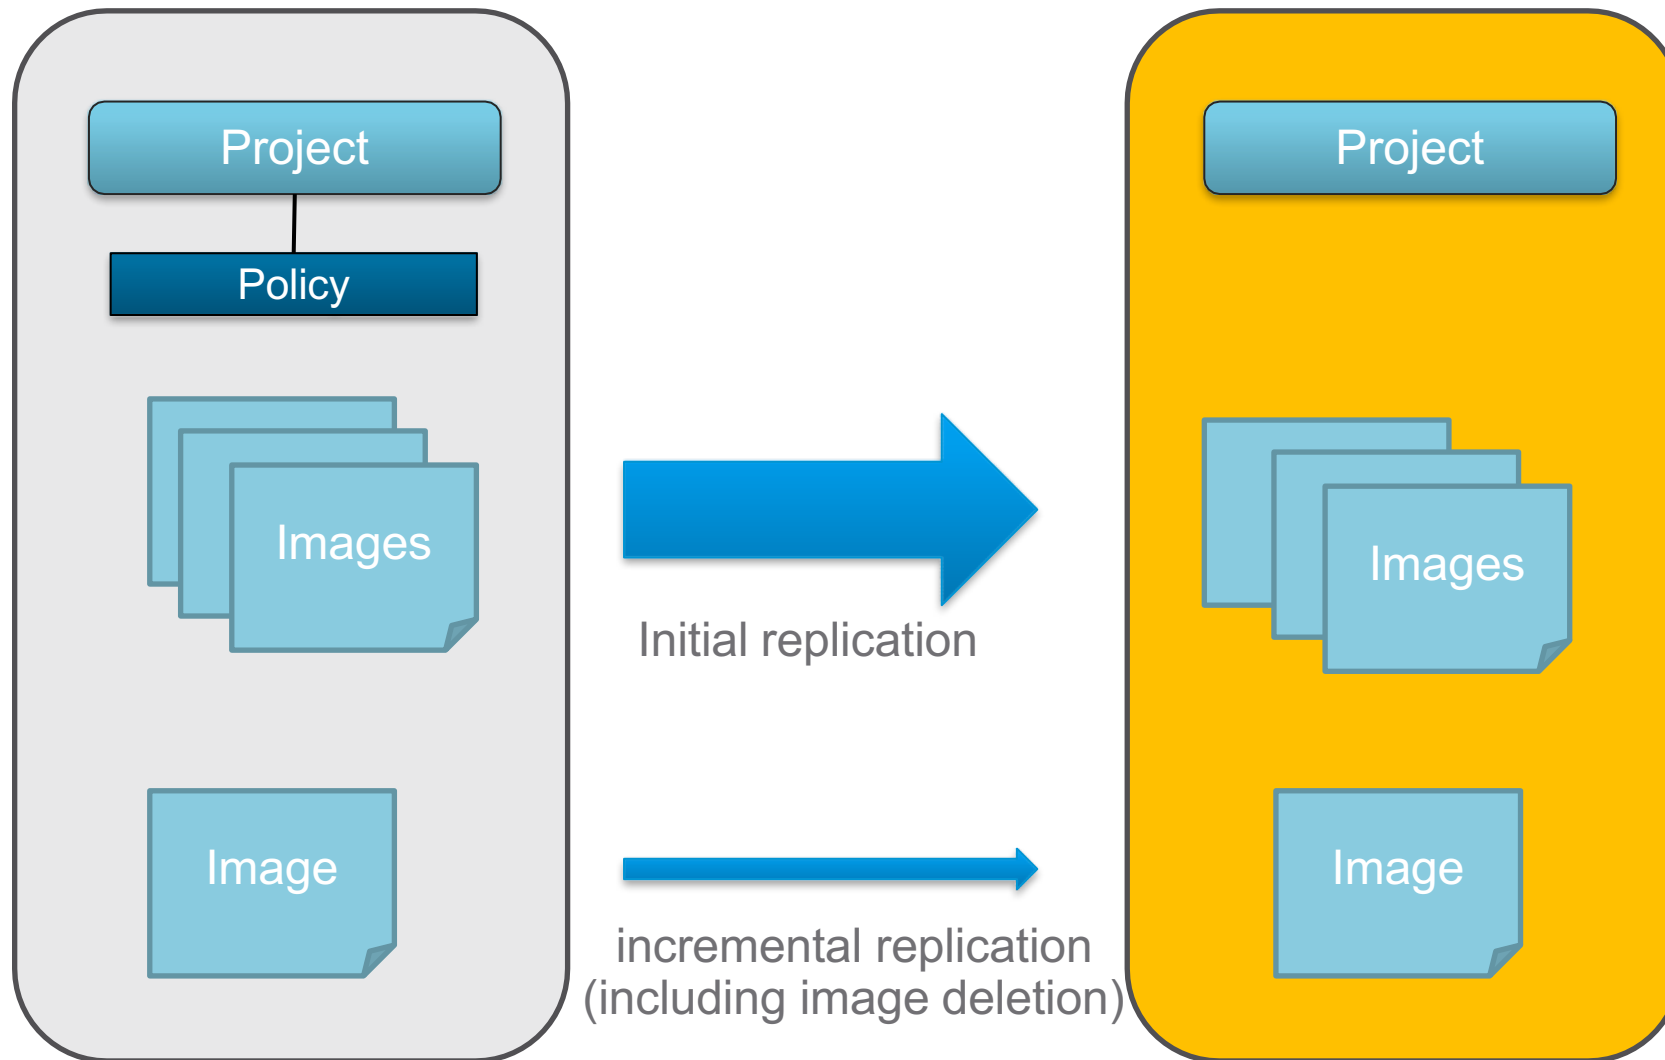


# Harbor Architecture





## Image replication (synchronization)



# Harbor

- 
- 1 Container Image Basics
  - 2 Project Harbor Introduction
  - 3 Consistency of Images
  - 4 Security
  - 5 Image Distribution
-

# Consistency of Container Images

- Container images are used throughout the life cycle of software development
  - Dev
  - Test
  - Staging
  - Production
- Consistency must be maintained
  - Version control
  - Issue tracking
  - Troubleshooting
  - Auditing

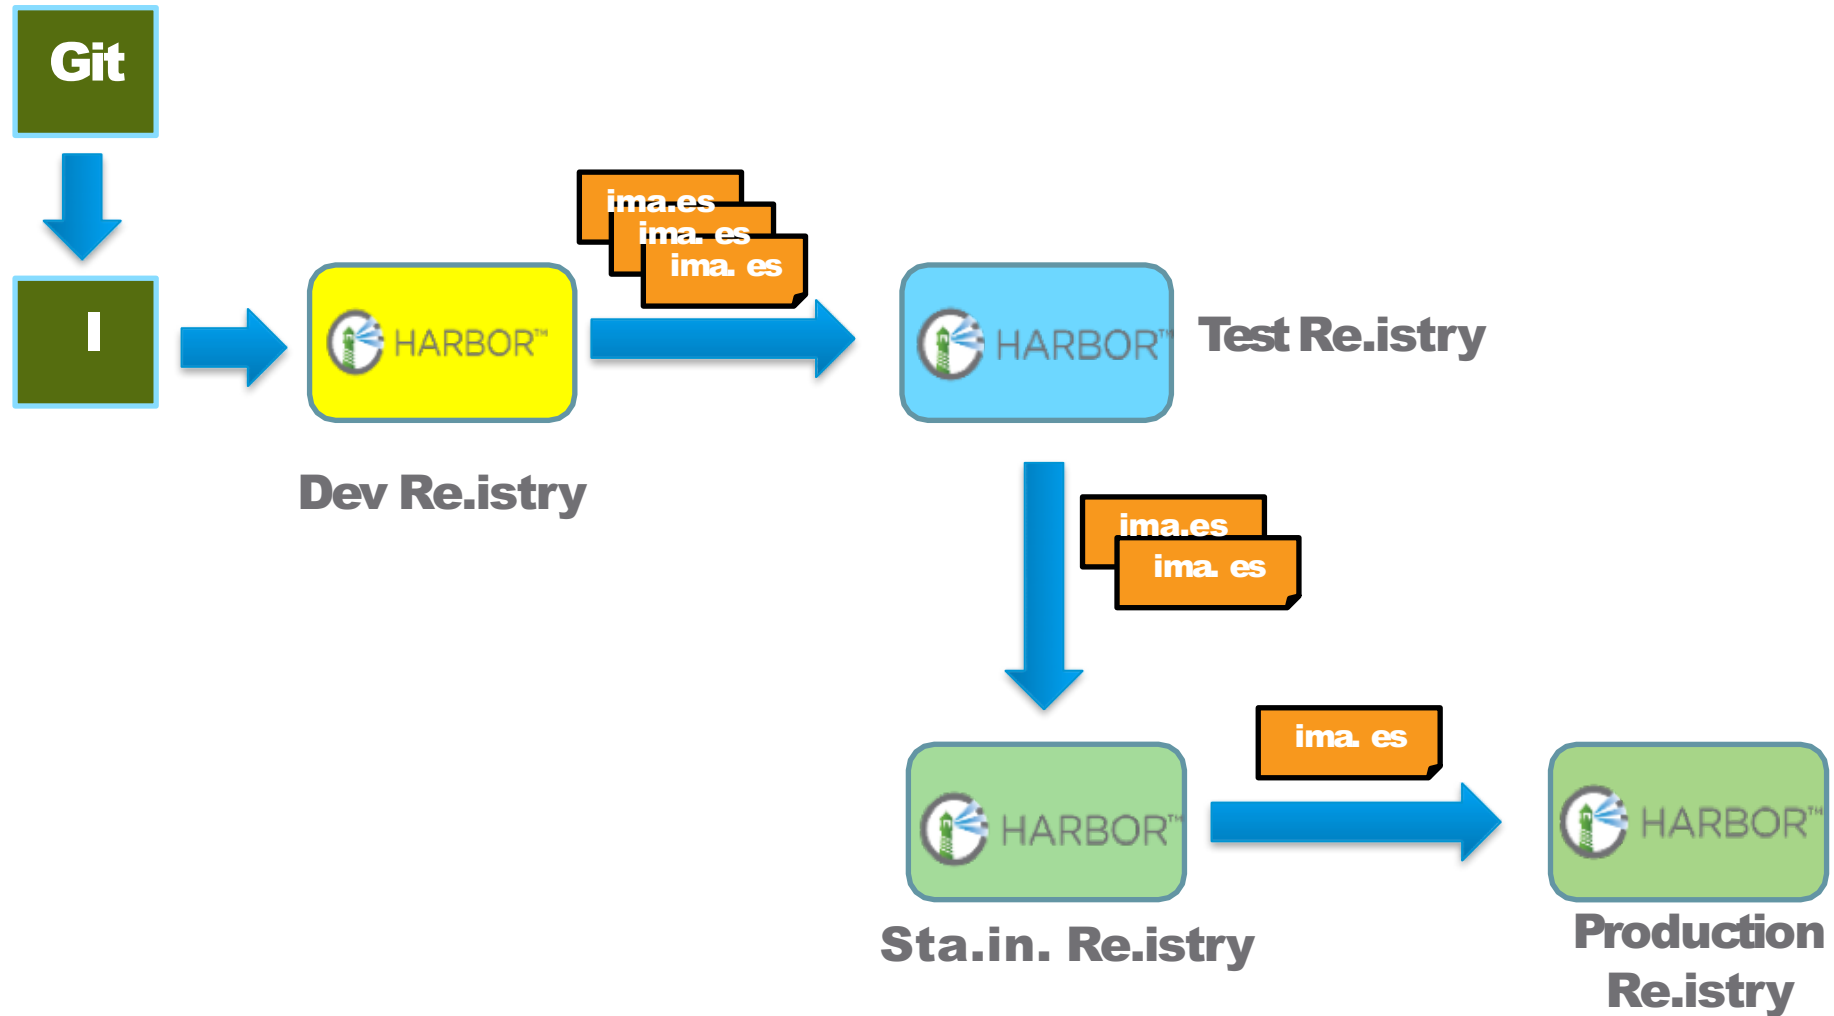
# Same Dockerfile Always Builds Same Image?

## Example:

```
FROM ubuntu
RUN apt-get install -y python
ADD app.jar /myapp/app.jar
```

- **Base image** `ubuntu:latest` could be changed between builds
- `ubuntu:14.04` could also be changed due to patching
- `apt-get (curl, wget..)` cannot guarantee always to install the same packages
- `ADD` depends on the build time environment to add files

# Shipping Images in Binary Format for Consistency



Ima.es are synchronized between environments  
by usin. Harbor re.istry.

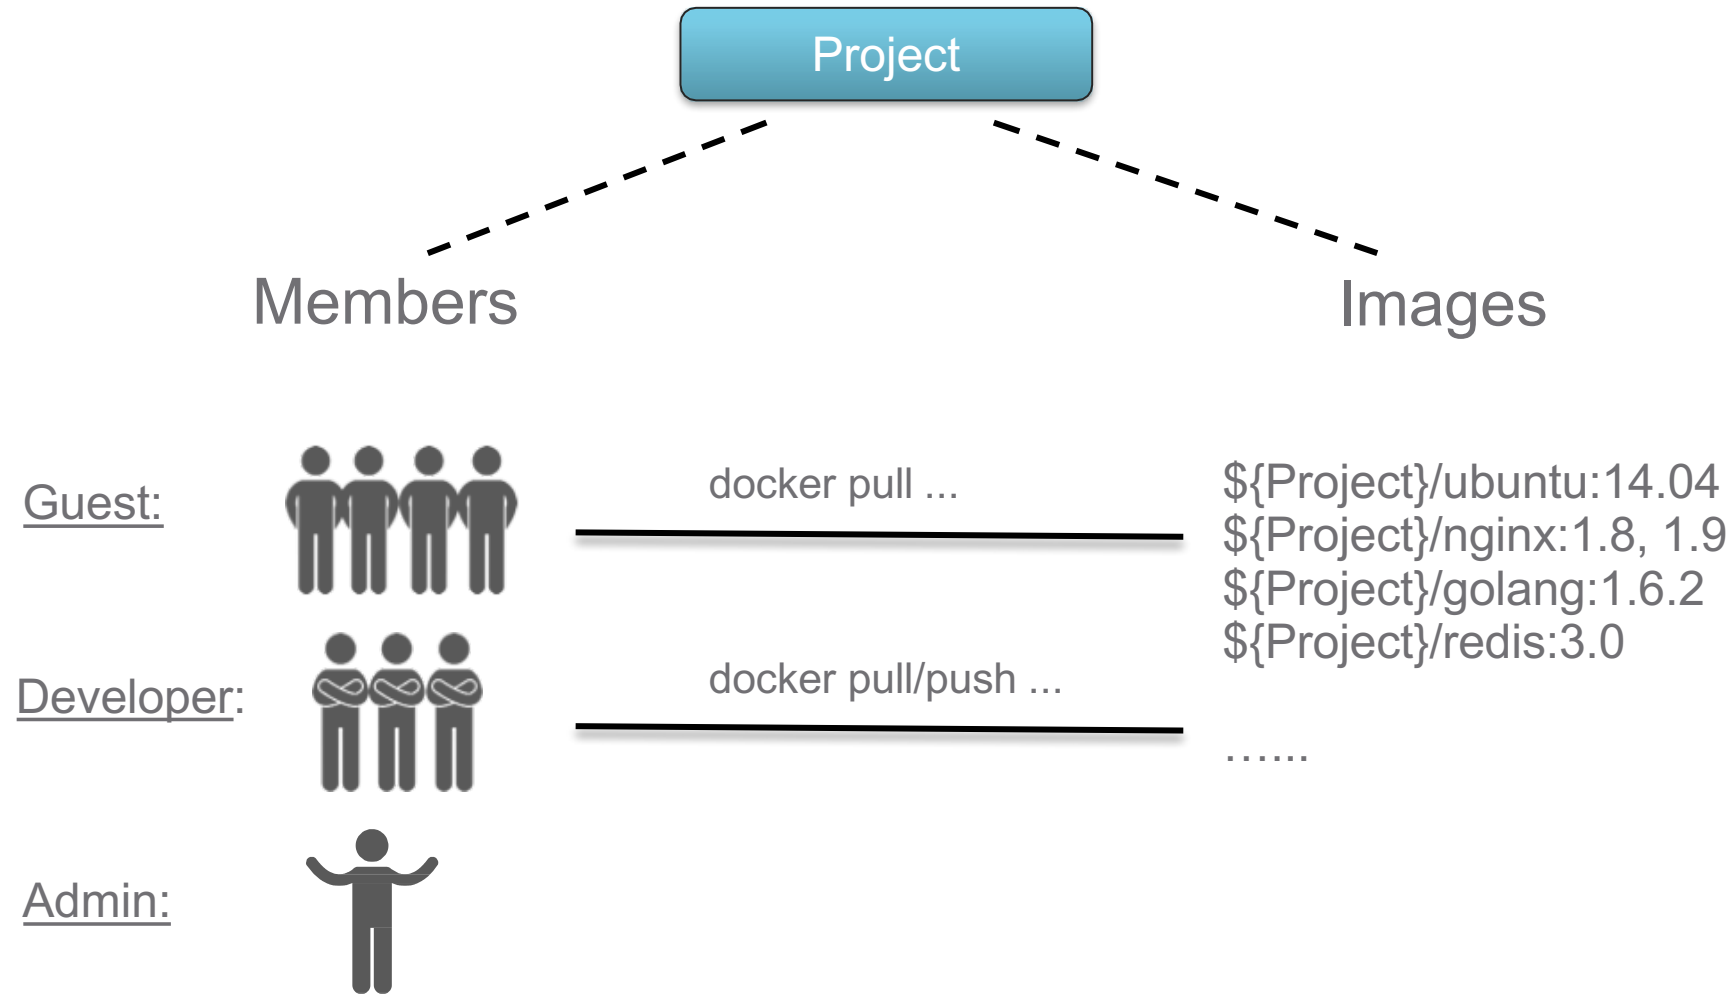
# Harbhor

- 
- 1 Container Image Basics
  - 2 Project Harbor Introduction
  - 3 Consistency of Images
  - 4 Security
  - 5 Image Distribution
-

# Access Control to Images

- Organizations often keep images within their own organizations
  - Intellectual property stays in organization
  - Efficiency: LAN vs WAN
- People with different roles should have different access
  - Developer – Read/Write
  - Tester – Read Only
- Different rules should be enforced in different environments
  - Dev/test env – many people can access
  - Production – a limited number of people can access
- Can be integrated with internal user management system
  - LDAP/Active Directory

# Example: Role Based Access Control in Harbor





## Other security considerations

- Enable content trust by installing Notary service
  - Image is signed by publisher's private key during pushing
  - Image is pulled using digest
- Perform vulnerability scanning
  - Identify images with vulnerabilities during pushing
  - Prevent images with vulnerabilities from being pulled
  - Regular scanning based on updated vulnerability database

# Harbor

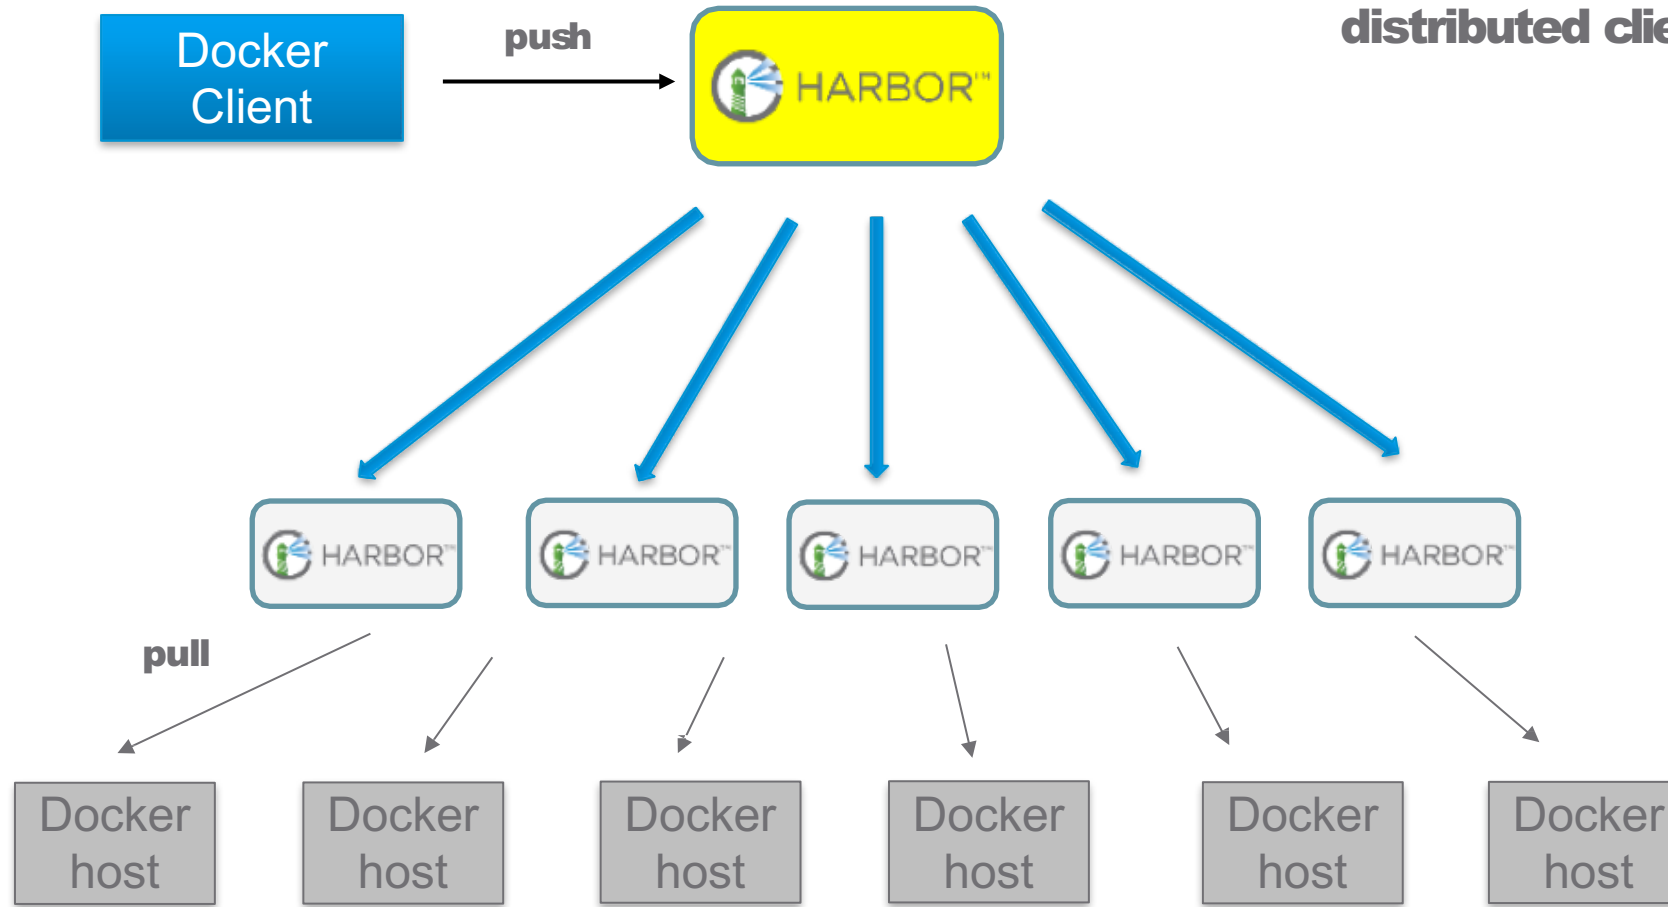
- 
- 1 Container Image Basics
  - 2 Project Harbor Introduction
  - 3 Consistency of Images
  - 4 Security
  - 5 Image Distribution
-

# Image Distribution

- Container images are usually distributed from a registry.
- Registry becomes the bottleneck for a large cluster of nodes
  - I/O
  - Network
- Scaling out an registry server
  - Multiple instances of registry sharing same storage
  - Multiple instances of independent registry sharing no storage

# Image Distribution via Master-Slave Replication

- Load balancin.
- Works well with .eo.raphically distributed clients



**Master – Slave model**

# Demo: Harbor Registry

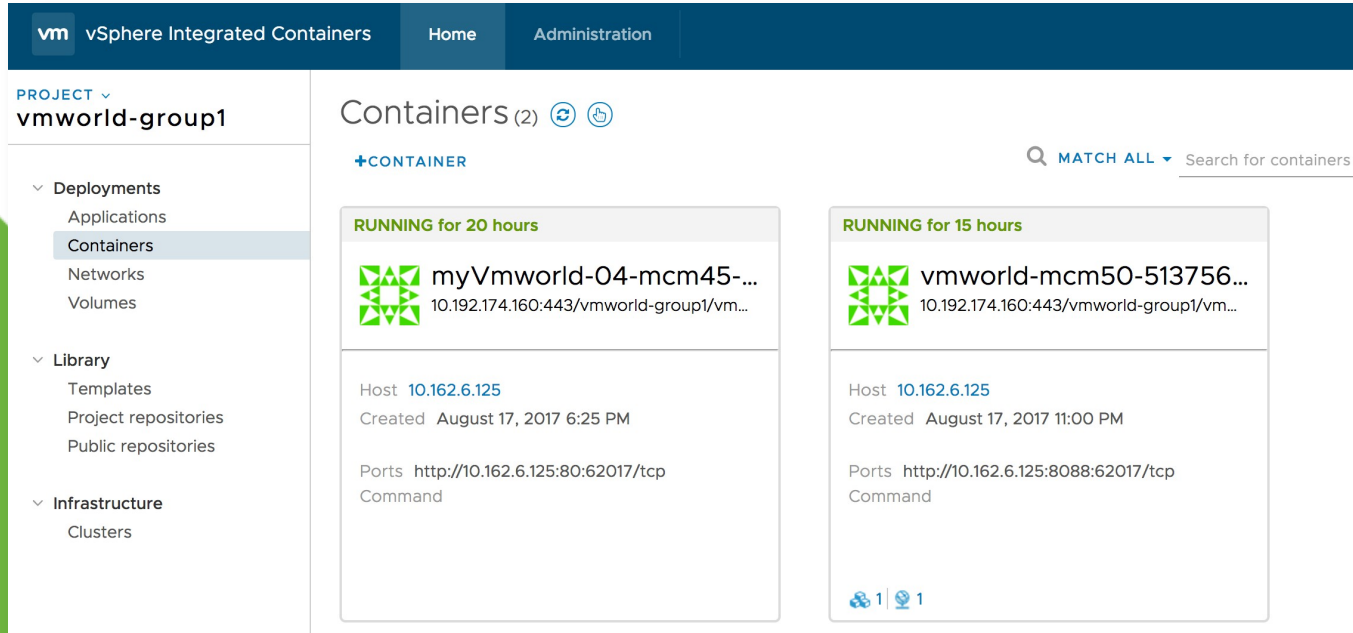
# Admiral Overview

## What is Admiral?

- A lightweight **Container Management Platform**
- Deploy and manage container based applications.
- Provide placement based on dynamic policy allocation.
- Container Management Layer in both vRA and VIC.
- Works with VIC engine or Docker host, Harbor registry.

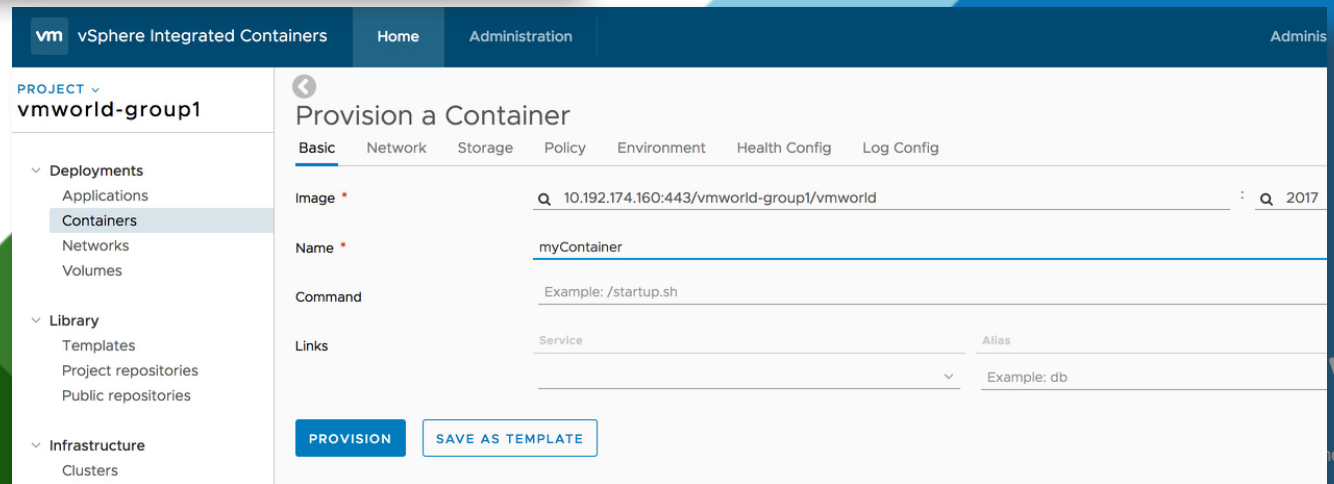
# Container Deployment

Provision container with the image from the specified registry



The screenshot shows the vSphere Integrated Containers web interface. The top navigation bar includes 'vm vSphere Integrated Containers', 'Home', and 'Administration'. The left sidebar shows the 'PROJECT' dropdown set to 'vmworld-group1' and a menu with 'Deployments' (Applications, Containers, Networks, Volumes), 'Library' (Templates, Project repositories, Public repositories), and 'Infrastructure' (Clusters). The main content area is titled 'Containers (2)' and shows two running containers:

- myVmworld-04-mcm45...**: RUNNING for 20 hours. Host: 10.162.6.125. Created: August 17, 2017 6:25 PM. Ports: http://10.162.6.125:80:62017/tcp. Command: Command.
- vmworld-mcm50-513756...**: RUNNING for 15 hours. Host: 10.162.6.125. Created: August 17, 2017 11:00 PM. Ports: http://10.162.6.125:8088:62017/tcp. Command: Command.



The screenshot shows the 'Provision a Container' form in the vSphere Integrated Containers web interface. The top navigation bar includes 'vm vSphere Integrated Containers', 'Home', 'Administration', and 'Adminis'. The left sidebar shows the 'PROJECT' dropdown set to 'vmworld-group1' and a menu with 'Deployments' (Applications, Containers, Networks, Volumes), 'Library' (Templates, Project repositories, Public repositories), and 'Infrastructure' (Clusters). The main content area is titled 'Provision a Container' and has tabs for 'Basic', 'Network', 'Storage', 'Policy', 'Environment', 'Health Config', and 'Log Config'. The 'Basic' tab is active, showing the following fields:

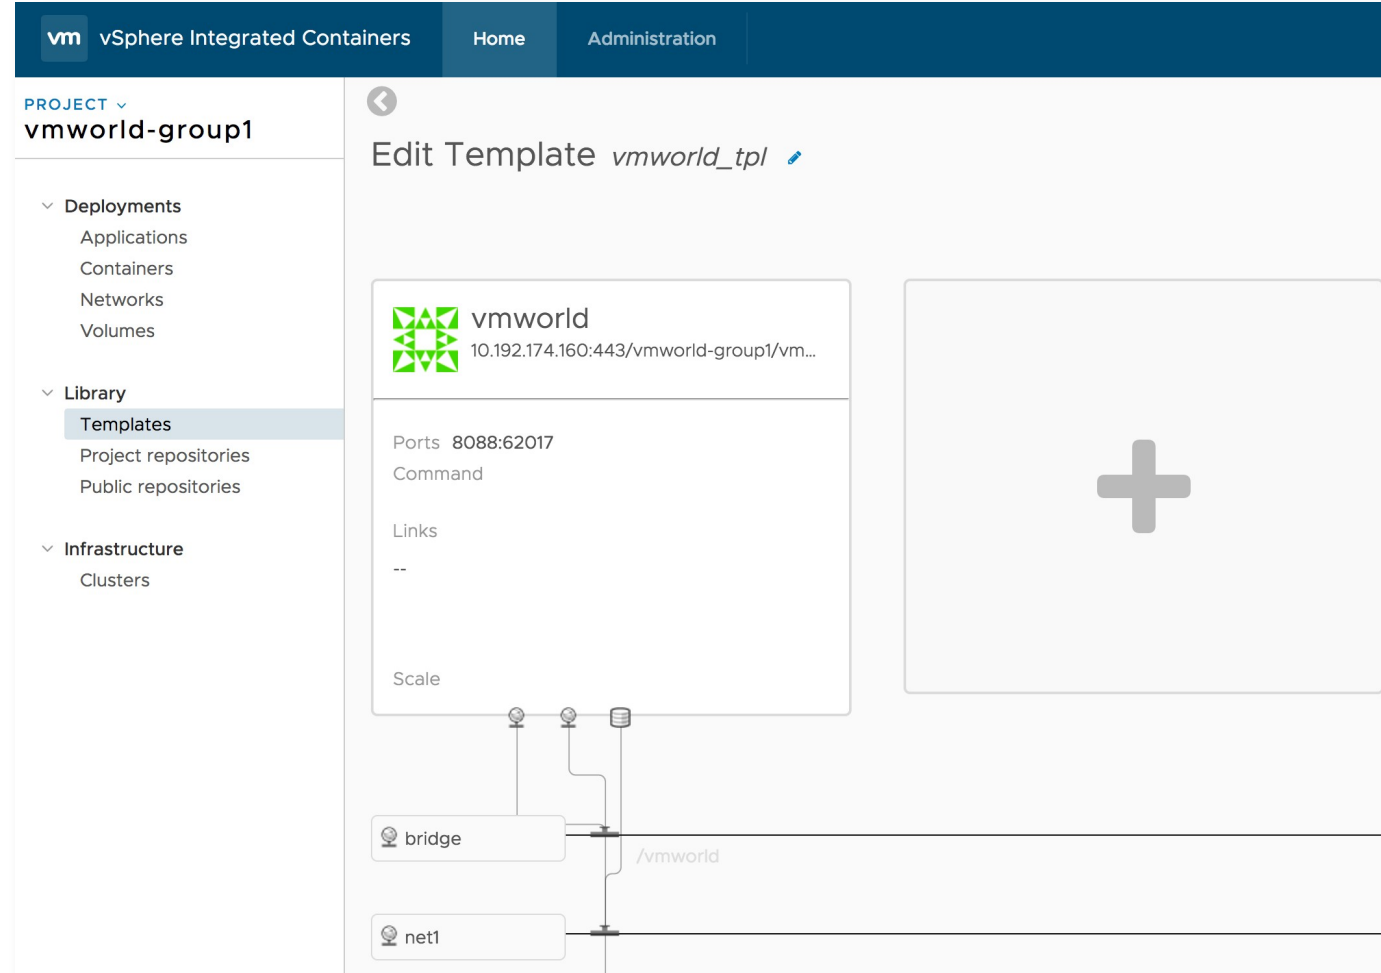
- Image \***: 10.192.174.160:443/vmworld-group1/vmworld
- Name \***: myContainer
- Command**: Example: /startup.sh
- Links**: Service (Example: db) and Alias (Example: db)

Buttons for 'PROVISION' and 'SAVE AS TEMPLATE' are at the bottom.



# Application Template

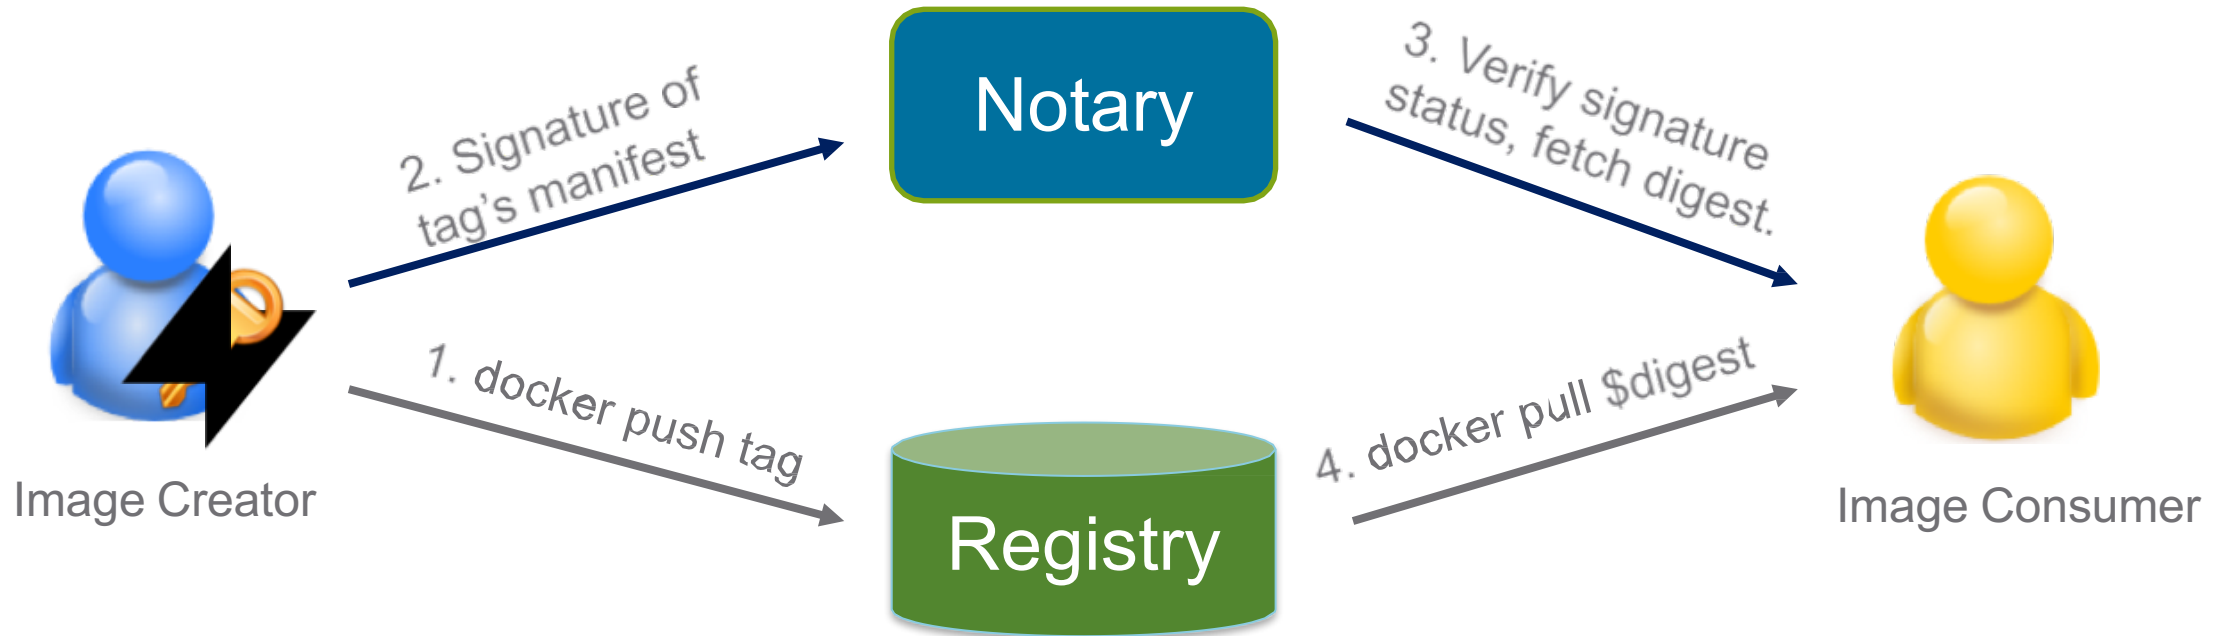
- Define your own application structure
  - Containers
  - Networks
  - Volumes
- Provision multiple containers at one time
- Compatible with **docker-compose** yaml file



The screenshot displays the VMware vSphere Integrated Containers management console. The top navigation bar includes 'vm vSphere Integrated Containers', 'Home', and 'Administration'. The left sidebar shows a project hierarchy for 'vmworld-group1' under 'PROJECT', with sections for Deployments (Applications, Containers, Networks, Volumes), Library (Templates, Project repositories, Public repositories), and Infrastructure (Clusters). The main content area is titled 'Edit Template vmworld\_tpl' and shows configuration details for a template named 'vmworld' with IP '10.192.174.160:443/vmworld-group1/vm...'. Fields for Ports (8088:62017), Command, Links, and Scale are visible. A network diagram at the bottom shows the template connected to a 'bridge' and 'net1' network.

# Practice on VIC engine, Harbor & Admiral

# Content trust for image provenance



# Vulnerability Scanning

- **Static analysis** of vulnerability by inspecting filesystem of container image and indexing features in database.
- **Rescanning** is needed only and only if new detectors are added.
- Update vulnerability data regularly
  - Debian Security Bug Tracker
  - Ubuntu CVE Tracker
  - Red Hat Security Data
  - Oracle Linux Security Data
  - Alpine SecDB